

Diabetes Prediction using Machine Learning

July 7, 2023

1 Importing the Dependencies

```
[1]: import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn import svm
from sklearn.metrics import accuracy_score
```

2 Data Collection and Analysis

```
[2]: df = pd.read_csv('diabetes.csv')
```

```
[3]: df
```

```
[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0

764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

```
[4]: #printing first 5 rows of the dataset
df.head()
```

```
[4]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0           6      148           72           35         0  33.6
1           1       85           66           29         0  26.6
2           8      183           64           0         0  23.3
3           1       89           66           23        94  28.1
4           0      137           40           35       168  43.1

      DiabetesPedigreeFunction  Age  Outcome
0                0.627      50         1
1                0.351      31         0
2                0.672      32         1
3                0.167      21         0
4                2.288      33         1
```

```
[5]: # no. of rows and columns in this dataset
df.shape
```

```
[5]: (768, 9)
```

```
[6]: #Getting the statistical measures of the data
df.describe()
```

```
[6]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  \
count  768.000000  768.000000    768.000000    768.000000  768.000000
mean    3.845052  120.894531    69.105469    20.536458    79.799479
std     3.369578   31.972618    19.355807    15.952218   115.244002
min     0.000000    0.000000     0.000000     0.000000     0.000000
25%     1.000000    99.000000    62.000000     0.000000     0.000000
50%     3.000000   117.000000    72.000000    23.000000    30.500000
75%     6.000000   140.250000    80.000000    32.000000   127.250000
max    17.000000   199.000000   122.000000    99.000000   846.000000

      BMI  DiabetesPedigreeFunction  Age  Outcome
count  768.000000    768.000000  768.000000  768.000000
mean    31.992578         0.471876   33.240885    0.348958
std     7.884160         0.331329   11.760232    0.476951
min     0.000000         0.078000   21.000000    0.000000
```

25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[7]: df['Outcome'].value_counts()
```

```
[7]: 0    500
      1    268
      Name: Outcome, dtype: int64
```

0 <- Non diabetic 1 <- Diabetic

```
[8]: df.groupby('Outcome').mean()
```

```
[8]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  \
Outcome
0      3.298000   109.980000      68.184000      19.664000   68.792000
1      4.865672   141.257463      70.824627      22.164179  100.335821

      BMI  DiabetesPedigreeFunction  Age
Outcome
0      30.304200                0.429734  31.190000
1      35.142537                0.550500  37.067164
```

```
[9]: # seperating data and labels
x = df.drop(columns = 'Outcome', axis= 1)
y= df['Outcome']
```

```
[10]: x
```

```
[10]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0              6      148           72           35         0  33.6
1              1       85           66           29         0  26.6
2              8      183           64            0         0  23.3
3              1       89           66           23        94  28.1
4              0      137           40           35       168  43.1
..          ...      ...           ...           ...      ...  ...
763           10      101           76           48       180  32.9
764              2      122           70           27         0  36.8
765              5      121           72           23       112  26.2
766              1      126           60            0         0  30.1
767              1       93           70           31         0  30.4

      DiabetesPedigreeFunction  Age
0                0.627      50
1                0.351      31
```

```

2          0.672  32
3          0.167  21
4          2.288  33
..          ...  ...
763        0.171  63
764        0.340  27
765        0.245  30
766        0.349  47
767        0.315  23

```

[768 rows x 8 columns]

```
[11]: y
```

```

[11]: 0      1
      1      0
      2      1
      3      0
      4      1
      ..
      763    0
      764    0
      765    0
      766    1
      767    0
Name: Outcome, Length: 768, dtype: int64

```

3 Data Standardization

```
[12]: scaler = StandardScaler()
```

```
[13]: scaler.fit(x)
```

```
[13]: StandardScaler()
```

```
[15]: standardized_data = scaler.transform(x)
```

```
[16]: print(standardized_data)
```

```

[[ 0.63994726  0.84832379  0.14964075 ...  0.20401277  0.46849198
   1.4259954 ]
 [-0.84488505 -1.12339636 -0.16054575 ... -0.68442195 -0.36506078
  -0.19067191]
 [ 1.23388019  1.94372388 -0.26394125 ... -1.10325546  0.60439732
  -0.10558415]
 ...

```

```
[ 0.3429808  0.00330087  0.14964075 ... -0.73518964 -0.68519336
 -0.27575966]
[-0.84488505  0.1597866  -0.47073225 ... -0.24020459 -0.37110101
  1.17073215]
[-0.84488505 -0.8730192   0.04624525 ... -0.20212881 -0.47378505
 -0.87137393]]
```

```
[17]: x = standardized_data
      y = df['Outcome']
```

```
[18]: x
```

```
[18]: array([[ 0.63994726,  0.84832379,  0.14964075, ...,  0.20401277,
             0.46849198,  1.4259954 ],
             [-0.84488505, -1.12339636, -0.16054575, ..., -0.68442195,
             -0.36506078, -0.19067191],
             [ 1.23388019,  1.94372388, -0.26394125, ..., -1.10325546,
             0.60439732, -0.10558415],
             ...,
             [ 0.3429808 ,  0.00330087,  0.14964075, ..., -0.73518964,
             -0.68519336, -0.27575966],
             [-0.84488505,  0.1597866 , -0.47073225, ..., -0.24020459,
             -0.37110101,  1.17073215],
             [-0.84488505, -0.8730192 ,  0.04624525, ..., -0.20212881,
             -0.47378505, -0.87137393]])
```

```
[19]: y
```

```
[19]: 0      1
      1      0
      2      1
      3      0
      4      1
      ..
     763     0
     764     0
     765     0
     766     1
     767     0
      Name: Outcome, Length: 768, dtype: int64
```

4 Train Test Split

```
[20]: x_train, x_test, y_train, y_test = train_test_split(x,y, test_size = 0.
      ↪2,stratify = y, random_state=2)
```

```
[21]: print(x.shape, x_train.shape, x_test.shape)
```

(768, 8) (614, 8) (154, 8)

5 Training the model

Support Vector Machine

```
[23]: classifier = svm.SVC(kernel = 'linear')
```

```
[24]: #training the support vector machine classifier  
classifier.fit(x_train, y_train)
```

```
[24]: SVC(kernel='linear')
```

6 Model Evaluation

Accuracy score

```
[25]: # accuracy score on the training data  
x_train_prediction = classifier.predict(x_train)  
training_data_accuracy = accuracy_score(x_train_prediction, y_train)
```

```
[26]: print('Accuracy score of the training data : ' , training_data_accuracy)
```

Accuracy score of the training data : 0.7866449511400652

```
[27]: # accuracy score on the test data  
x_test_prediction = classifier.predict(x_test)  
test_data_accuracy = accuracy_score(x_test_prediction, y_test)
```

```
[28]: print('Accuracy score of the test data : ' , test_data_accuracy)
```

Accuracy score of the test data : 0.7727272727272727

7 Building a predictive system

```
[30]: input_data = (4,110,92,0,0,37.6,0.191,30)
```

```
[32]: input_data = (4,110,92,0,0,37.6,0.191,30)  
# Changing the input data into numpy array  
input_data_as_numpy_array = np.asarray(input_data)  
  
#reshape the array as we are predicting for one instance  
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)
```

```

#standardize the input data
std_data = scaler.transform(input_data_reshaped)
print(std_data)

prediction = classifier.predict(std_data)
print(prediction)

if (prediction[0]==0):
    print('The person is not diabetic')
else:
    print('The person is diabetic')

```

```

[[ 0.04601433 -0.34096773  1.18359575 -1.28821221 -0.69289057  0.71168975
 -0.84827977 -0.27575966]]

```

```
[0]
```

The person is not diabetic

C:\Users\DELL\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names

```
warnings.warn(
```